

Problem Definition:

To examine customer data from IBM Sample Data Sets with the aim of building and comparing several customer churn prediction models. The data contain information such as 'customerID', 'gender', 'SeniorCitizen', 'Partner', etc. Dataset has 7043 instances and 21 attributes containing a blend of categorical and numerical values.

Data Analysis:

Before moving to model building and evaluation phase, data analysis helps in understanding the data and deriving insights about dataset. Data analysis requires cleaning, transforming and modelling of data to identify useful information from data and taking decision on basis of derived result.

First, import the required libraries

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
```

Now, Load the data into pandas DataFrame using read_csv function.

```
df=pd.read_csv("""Customer_churn_analysis.csv""")
```

```
df.shape
```

Shape attribute provides the information about the dataset containing 7043 rows and 21 columns

Info function provides concise summary of the DataFrame including the datatype

```
df.info ()
```

customerID	7043	non-null	object
gender	7043	non-null	object
SeniorCitizen	7043	non-null	int64
Partner	7043	non-null	object
Dependents	7043	non-null	object
tenure	7043	non-null	int64
PhoneService	7043	non-null	object
MultipleLines	7043	non-null	object
InternetService	7043	non-null	object
OnlineSecurity	7043	non-null	object
OnlineBackup	7043	non-null	object
DeviceProtection	7043	non-null	object
TechSupport	7043	non-null	object
StreamingTV	7043	non-null	object
StreamingMovies	7043	non-null	object

```
Contract          7043 non-null object
PaperlessBilling  7043 non-null object
PaymentMethod     7043 non-null object
MonthlyCharges    7043 non-null float64
TotalCharges      7043 non-null object
Churn             7043 non-null object
```

It can be observed that no missing values can be identified with attributes a combination of numerical and categorical datatypes.

Handling numerical columns:

```
df.describe()
```

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

Describe function give us summary of statistics regarding the DataFrame columns. It displays the mean, std and IQR values for numeric columns.


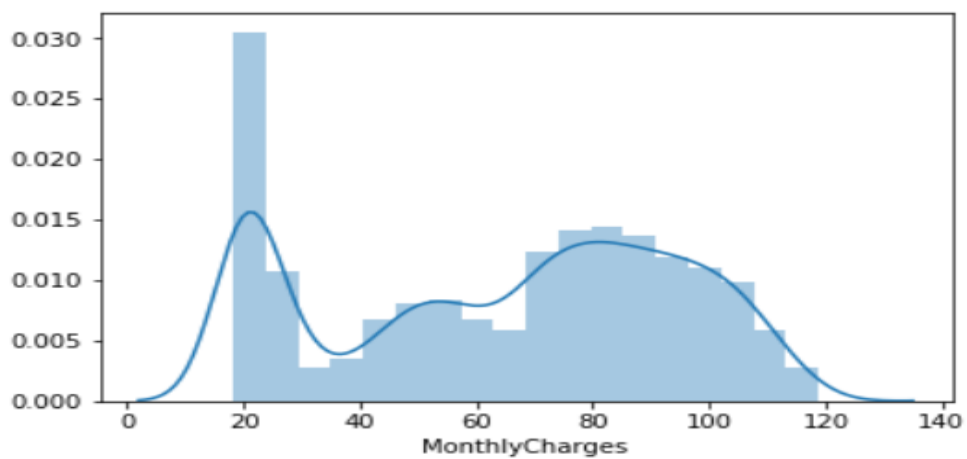
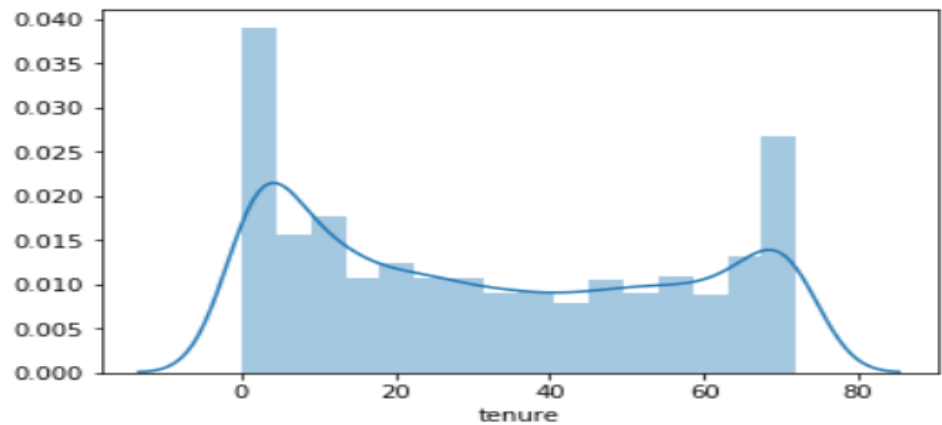
It can be identified that SeniorCitizen has median as 0 and mean as 0.1621. Similarly mean for tenure is 32.37 and for MonthlyCharges it is 64.76.

```
df["TotalCharges"].value_counts()
df["MonthlyCharges"].value_counts()
df["tenure"].value_counts()
```

Using the value_counts() function it can easily be noticed that "Totalcharges" have datatype as object but values are float type so need to convert datatype

Visualization of numerical columns


```
sns.distplot(df["tenure"])  
sns.distplot(df["MonthlyCharges"])  
sns.distplot(df["TotalCharges"])
```



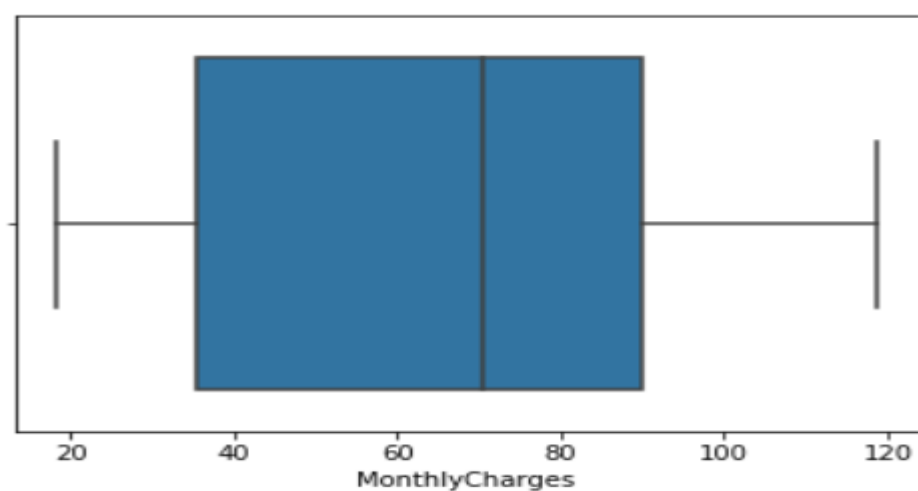
Using the distribution plot variation in data distribution can be identified. Also it helps us understanding the skewness in the data. From distribution plot it can be identified that "tenure", "MonthlyCharges", "TotalCharges" are not equally distributed and are skewed. Skewness can be reduced using Power Transform, Log Transform, Square root Transform, Box-Cox Transform.

Along with skewness, outliers must also be considered so that model accuracy is not affected. To check the outliers, boxplot can be used and outliers can be removed using zscore.

```
sns.boxplot(x="tenure", data=df)
```



```
sns.boxplot(x="MonthlyCharges", data=df)
```




Looking at the boxplots, it can be identified that “tenure”, “MonthlyCharges”, do not have outlier values.


Handling Categorical columns

Categorical columns can be visualized using countplot as it provides the counts of observations in each categorical bin using bars.


```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="gender",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="Partner",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="Dependents",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="PhoneService",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="MultipleLines",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="InternetService",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="OnlineSecurity",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="OnlineBackup",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="DeviceProtection",data=df,hue="Churn")
```


```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="TechSupport",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="StreamingTV",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")  
sns.countplot(x="StreamingMovies",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")  
sns.countplot(x="Contract",data=df,hue="Churn")
```




```
plt.figure(figsize=(15,10),facecolor="white")  
sns.countplot(x="PaperlessBilling",data=df,hue="Churn")
```



```
plt.figure(figsize=(15,10),facecolor="white")  
sns.countplot(x="PaymentMethod",data=df,hue="Churn")
```



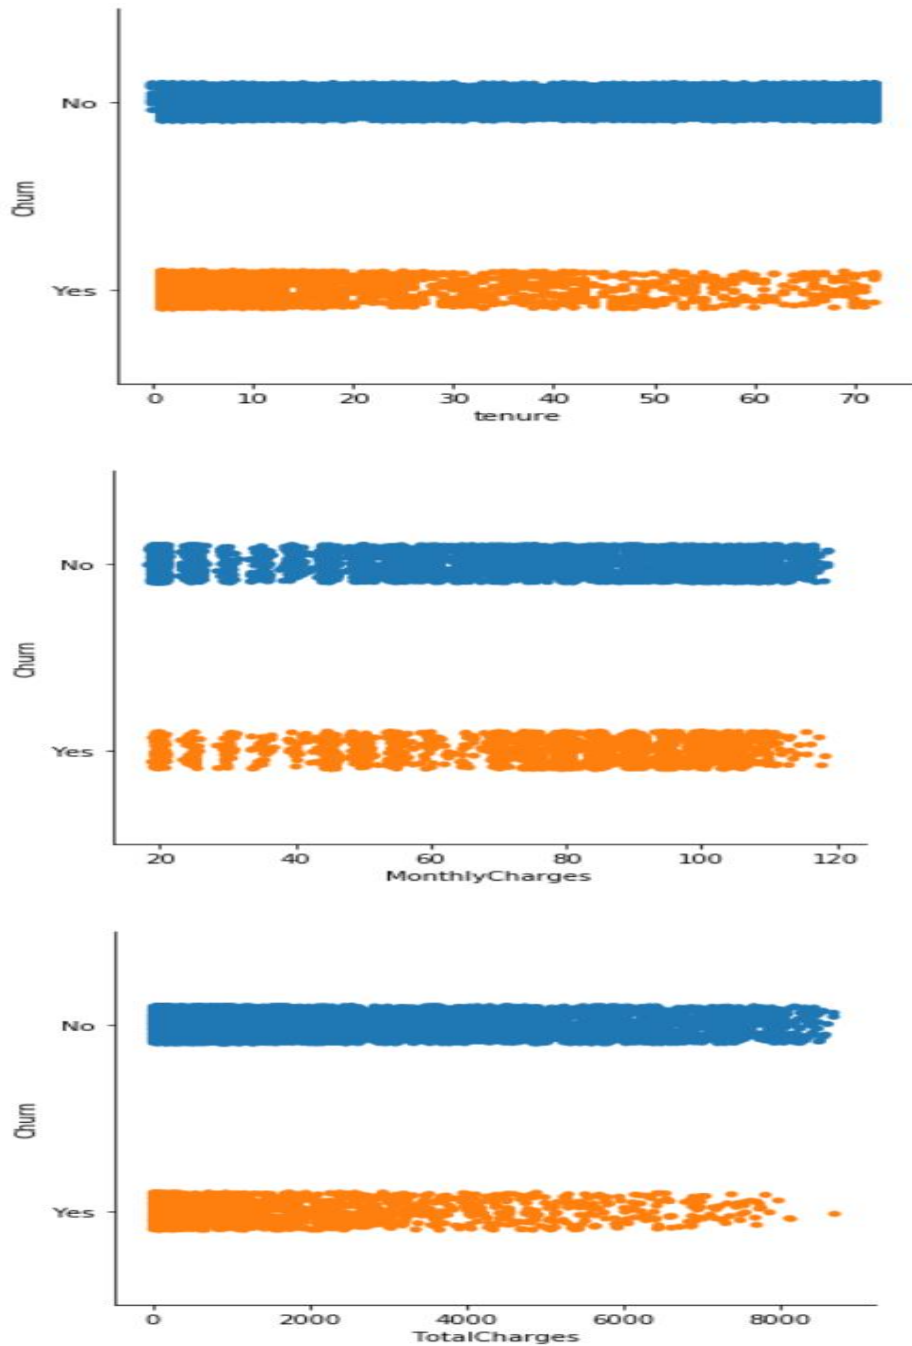
```
plt.figure(figsize=(15,10),facecolor="white")
sns.countplot(x="SeniorCitizen",data=df,hue="Churn")
```



Using the countplot it can be noticed that columns "PhoneService", "MultipleLines", "SeniorCitizen", have class imbalance problem which can be resolved using undersampling or oversampling technique.

Catplot

```
sns.catplot(x="tenure", y="Churn", data=df)
sns.catplot(x="MonthlyCharges", y="Churn", data=df)
sns.catplot(x="TotalCharges", y="Churn", data=df)
```



Correlation

Correlation is really important to identify which columns are affecting the target column positively or negatively. Through correlation it can also be identified that which columns are strongly related to target.

```
df.corr()
```

For plotting correlation heatmap can be used to understand the correlation more easily.

```
import matplotlib.pyplot as plt
plt.figure(figsize=(30,15))
sns.heatmap(df.corr(),annot=True,linewidths=0.1,linecolor="black",fmt=".2f")
```

Through correlation it can be identified that **tenure**, **Contract** have strong negative correlation with **churn**.

EDA Concluding Remark:

- Dataset contains 7043 rows and 21 columns with no missing values
- Dataset comprise of column of categorical as well as numerical type.
- It can be observed that no missing values can be identified
- It can easily be noticed that “Totalcharges” have datatype as object but values are continuous so need to convert datatype.
- “customerID” will not add any significance to the model building process so it must be dropped.
- It can be identified that “tenure”, “MonthlyCharges”, “TotalCharges” are not equally distributed and are skewed.
- It can be identified that “tenure”, “MonthlyCharges”, do not have any outlier values.
- it can be noticed that columns "PhoneService", "MultipleLines", "SeniorCitizen", have class imbalance problem which can be resolved using undersampling or oversampling technique.
- It can be identified that **tenure**, **Contract** have strong negative correlation with **churn**.
- MinMaxScaler can be used for the normalization of data.

Pre-Processing:

Converting the datatype of “Totalcharges” from object to float type

```
df["TotalCharges"] = df["TotalCharges"].str.strip()
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"])
```

Dropping the column "customerID"

```
df.drop(columns=["customerID"], inplace=True)
```

To handle the numerical and categorical attributes differently. Numerical attributes need to be scaled, whereas for categorical columns missing values need to be filled and then encode the categorical values into numerical values.

We have lot of features which are categorical and for building a model, categorical column need to be encoded. OrdinalEncoder can be used to encode the categorical columns

```
from sklearn.preprocessing import OrdinalEncoder
enc=OrdinalEncoder()
for i in df.columns:
    if df[i].dtypes=="object":
        df[i]=enc.fit_transform(df[i].values.reshape(-1,1))
```

Before proceeding to model fitting, data transformation must be performed. Properly formatted and validated data improves data quality. Data transformation allow scaling each feature to a given range. For the given dataset MinMaxScaler can be used which scales all the data features in the range [0,1] or else in the range [-1,1].

Here, MinMaxScaler will be used to normalize the data column by column.

```
from sklearn.preprocessing import MinMaxScaler
scale=MinMaxScaler()
df["tenure"]=scale.fit_transform(df["tenure"].values.reshape(-1,1))
```

```
from sklearn.preprocessing import MinMaxScaler
scale=MinMaxScaler()
df["MonthlyCharges"]=scale.fit_transform(df["MonthlyCharges"].values.reshape(-1,1))
```

```
from sklearn.preprocessing import MinMaxScaler
scale=MinMaxScaler()
df["TotalCharges"]=scale.fit_transform(df["TotalCharges"].values.reshape(-1,1))
```

Building Machine Learning Models:

Model Building involves splitting of original dataset into input(x) and output (y) columns.

```
y=df["Churn"]
x=df.drop("Churn",axis=1)
```

x and y will be used to divide the dataset into training dataset and testing dataset.

Training dataset can be used for model fitting and statistical method can be used to estimate the accuracy of the model using the unseen data. Testing set can be used for evaluating the model accuracy.

Loaded dataset will be splitted into two, 70% of which will be used to train and 30% will be used to test the model.

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, roc_auc_score
from sklearn.metrics import confusion_matrix, roc_auc_score
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings("ignore")
```

Random state allows to generate a reproducible split. Scikit-learn use random permutations to generate the splits. Random state provided is a seed to generate random number and also ensure that generated numbers are in same order.

Maximum accuracy_score can be used to identify the best random state from 1 to 200 and will proceed with same random state to fit the model using different algorithms.

```
Maxr2=0
BestRs=0
for i in range(1,200):
    xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=.30,random_state=i)
    lr=LogisticRegression()
    lr.fit(xtrain,ytrain)
    pred=lr.predict(xtest)
    accu=accuracy_score(pred,ytest)
    print(accu)
    roc_score=roc_auc_score(pred,ytest)
    if roc_score>Maxr2:
        Maxr2=roc_score
        BestRs=i
print("with random state as",BestRs,"max accuracy is",Maxr2)
```

```
with random state as 80 max accuracy is 0.7773657420152327
```

So maximum roc_auc_score is .7773 which mean 77.73% accuracy at random state 80.

Following are the metrics we'll use to evaluate our predictive accuracy:

- Sensitivity = True Positive Rate (TP/TP+FN) – It says, 'out of all the positive (majority class) values, how many have been predicted correctly'.
- Specificity = True Negative Rate (TN/TN +FP) – It says, 'out of all the negative (minority class) values, how many have been predicted correctly'.
- Precision = (TP/TP+FP)
- Recall = Sensitivity
- F score = $2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$ – It is the harmonic mean of precision and recall. It is used to compare several models side-by-side. Higher the better.

Logistic Regression Model

So we will fit logistic regression model with random state=80

```
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=.30,random_state=80)
lr=LogisticRegression()
lr.fit(xtrain,ytrain)
pred=lr.predict(xtest)
accu=accuracy_score(pred,ytest)
print(classification_report(pred,ytest))
print(accu)
print(confusion_matrix(pred,ytest))
print(roc_auc_score(pred,ytest))
```

	precision	recall	f1-score	support
0.0	0.92	0.84	0.88	1669
1.0	0.55	0.71	0.62	444
accuracy			0.82	2113
macro avg	0.73	0.78	0.75	2113
weighted avg	0.84	0.82	0.82	2113

```
accuracy score = 0.8154283009938476

confusion matrix = [[1407  262]
                    [ 128  316]]

roc_auc_score = 0.7773657420152327
```

Outcome of the logistic regression model gave us accuracy score (0.815) and roc_auc_score as .777

Random Forest Model

So, fitting a Random Forest model with random state 80

```
from sklearn.ensemble import RandomForestClassifier

RFC=RandomForestClassifier()
RFC.fit(xtrain,ytrain)
pred=RFC.predict(xtest)
accu=accuracy_score(pred,ytest)
print(classification_report(pred,ytest))
print(accu)
print(confusion_matrix(pred,ytest))
print(roc_auc_score(pred,ytest))
```

	precision	recall	f1-score	support
0.0	0.91	0.82	0.86	1692
1.0	0.48	0.67	0.56	421
accuracy			0.79	2113
macro avg	0.70	0.74	0.71	2113
weighted avg	0.82	0.79	0.80	2113

```
accuracy score = 0.792238523426408

confusion matrix = [[1394  298]
                    [ 141  280]]

roc_auc_score = 0.7444801019749219
```

Outcome of the Random forest classifier model gave us accuracy score (0.792) and roc_auc_score is .744 .

Decision Tree Model

So, fitting a Decision Tree model with random state 80

```
from sklearn.tree import DecisionTreeClassifier

dtc=DecisionTreeClassifier()
dtc.fit(xtrain,ytrain)
pred=dtc.predict(xtest)
accu=accuracy_score(pred,ytest)
print(classification_report(pred,ytest))
print(accu)
print(confusion_matrix(pred,ytest))
```

```
print(roc_auc_score(pred,ytest))
```

```

              precision    recall  f1-score   support

    0.0         0.81         0.81         0.81        1521
    1.0         0.51         0.49         0.50         592

 accuracy
macro avg         0.66         0.65         0.65        2113
weighted avg         0.72         0.72         0.72        2113

accuracy score = 0.7236157122574538

confusion matrix = [[1236  285]
                   [ 299  293]]

roc_auc_score = 0.653777853297084

```

Outcome of the Decision Tree Classifier model gave us accuracy score (0.723) and roc_auc_score as 0.653 .

Support Vector Machine Model

So, fitting a support vector machine classifier model with random state 80

```

from sklearn.svm import SVC

svc=SVC()
svc.fit(xtrain,ytrain)
pred=svc.predict(xtest)
accu=accuracy_score(pred,ytest)
print(classification_report(pred,ytest))
print(accu)
print(confusion_matrix(pred,ytest))
print(roc_auc_score(pred,ytest))

```

```

              precision    recall  f1-score   support

    0.0         0.93         0.82         0.87        1735
    1.0         0.47         0.71         0.56         378

 accuracy
macro avg         0.70         0.77         0.72        2113
weighted avg         0.85         0.80         0.82        2113

accuracy score = 0.8021769995267393

confusion matrix = [[1426  309]
                   [ 109  269]]

```

0.7667711144656388

Outcome of the support vector classifier model gave us accuracy score (0.802) and auc_roc_score (0.766)

Cross validation

Cross Validation is a technique for assessing how the statistical analysis generalises to an independent dataset. It is a technique for evaluating machine learning models by training several models on subsets of the available input data and evaluating them on the complementary subset of the data

For a k fold cross-validation, k refers to the number of groups a given data sample is to be split into and the procedure is often called k-fold cross-validation. When a specific value for k is chosen, then it may be used such as k=5

```
from sklearn.model_selection import cross_val_score
model=[lr,dtc,RFC,svc]

for model in model:
    print(cross_val_score(model,x,y,cv=5,scoring="roc_auc").mean())
```

0.8436054857675283

0.652881963774847

0.8216564299596982

0.804040467355892

After comparing the roc_auc_score and cross_val_score for logistic regression, decision tree classifier, random forest classifier and support vector classifier, it can be noticed that difference in roc_auc_score and cross_val_score for SVC (Support Vector Classifier) is lowest implying that Support Vector Classifier model is more accurate for prediction.

So we will proceed with Support Vector Classifier and will try to tune hyperparameter using GridSearchCV

Hyperparameter tuning

Hyperparameter is a parameter which control the learning rate. Hyperparameters tuning is crucial as they control the overall behavior of a machine learning model. Every machine learning models will have different hyperparameters that can be set. A hyperparameter is a parameter whose value is set before the learning process begins

Tuning the parameters of the Support Vector Classifier in order to obtain the best possible parameters for model building.

```
from sklearn.model_selection import GridSearchCV
```

```
parameter={"kernel":["linear","poly","rbf","sigmoid"],"gamma":["scale","auto"]}
grid=GridSearchCV(estimator=svc,param_grid=parameter,cv=5)
grid.fit(xtrain,ytrain)
print(grid.best_score_)
print(grid.best_params_)
```

```
Best score = 0.79026369168357
```

```
Best parameter = {'gamma': 'scale', 'kernel': 'linear'}
```

Using gridsearchcv best parameters are identified which can used to develop final Support Vector Classifier model. Best score which is obtained using the best parameters is .7902 .

So we will be developing final model using the best parameters.

```
{'gamma': 'scale', 'kernel': 'linear'}
```

Final Model

```
finalsvc=SVC(kernel='linear',gamma="scale")
finalsvc.fit(xtrain,ytrain)
pred=finalsvc.predict(xtest)
print("classification_report :",classification_report(pred,ytest))
print("accuracy_score :",accuracy_score(pred,ytest))
print("confusion_matrix :",confusion_matrix(pred,ytest))
print("roc_auc_score :",roc_auc_score(pred,ytest))
```

	precision	recall	f1-score	support
0.0	0.90	0.84	0.87	1644
1.0	0.54	0.67	0.60	469
accuracy			0.80	2113
macro avg	0.72	0.75	0.73	2113
weighted avg	0.82	0.80	0.81	2113

```
accuracy_score : 0.8017037387600567
```

```
confusion_matrix : [[1380  264]
                   [ 155  314]]
```

```
roc_auc_score : 0.7544628266384449
```

Final model which is finalsvc has an accuracy_score of .8017 and roc_auc_score of .7544. Now we will be saving the final model and 75.44% accuracy means that it can be predicted that a customer will stop doing business with company or not with 75.44% accuracy.

Saving the final model

Saving the model using joblib library

```
import joblib
joblib.dump(finalsvc,"customer_churn.obj")
```

Concluding Remarks:

Here by using sklearn, we have built a preliminary machine learning tool and we have used the roc_auc_score to select the support vector classifier model as the best model for predicting that a customer will stop doing business with company or not. Our final model has an accuracy of 75.44% using the parameter tuning. We can further try to improve the model prediction by working on preprocessing and roc_auc_score. Of course, there is always tools and analysis you can do further in order to make it more accurate, and better to use.